

Distributed Machine Learning Toolkit (DMTK)

Microsoft Research Asia


2015.11

Overview

- Introduction to DMTK
- Using Binary Tools in DMTK
- Developing Your Own Algorithms on DMTK
- Resource and Support

DMTK

- DMTK (current release) includes
 - Distributed machine learning framework (code name: **Multiverso**)
 - Parameter server + client SDK for developing distributed machine learning algorithms
 - Two distributed algorithms implemented on Multiverso
 - LightLDA
 - Distributed trainer for LightLDA^[1] topic model
 - Distribute Word Embedding(DWE)
 - Distributed trainer for word2vec^[2] (abbr. as DWE) and multi-sense word embedding^[3] (abbr. as DMWE).

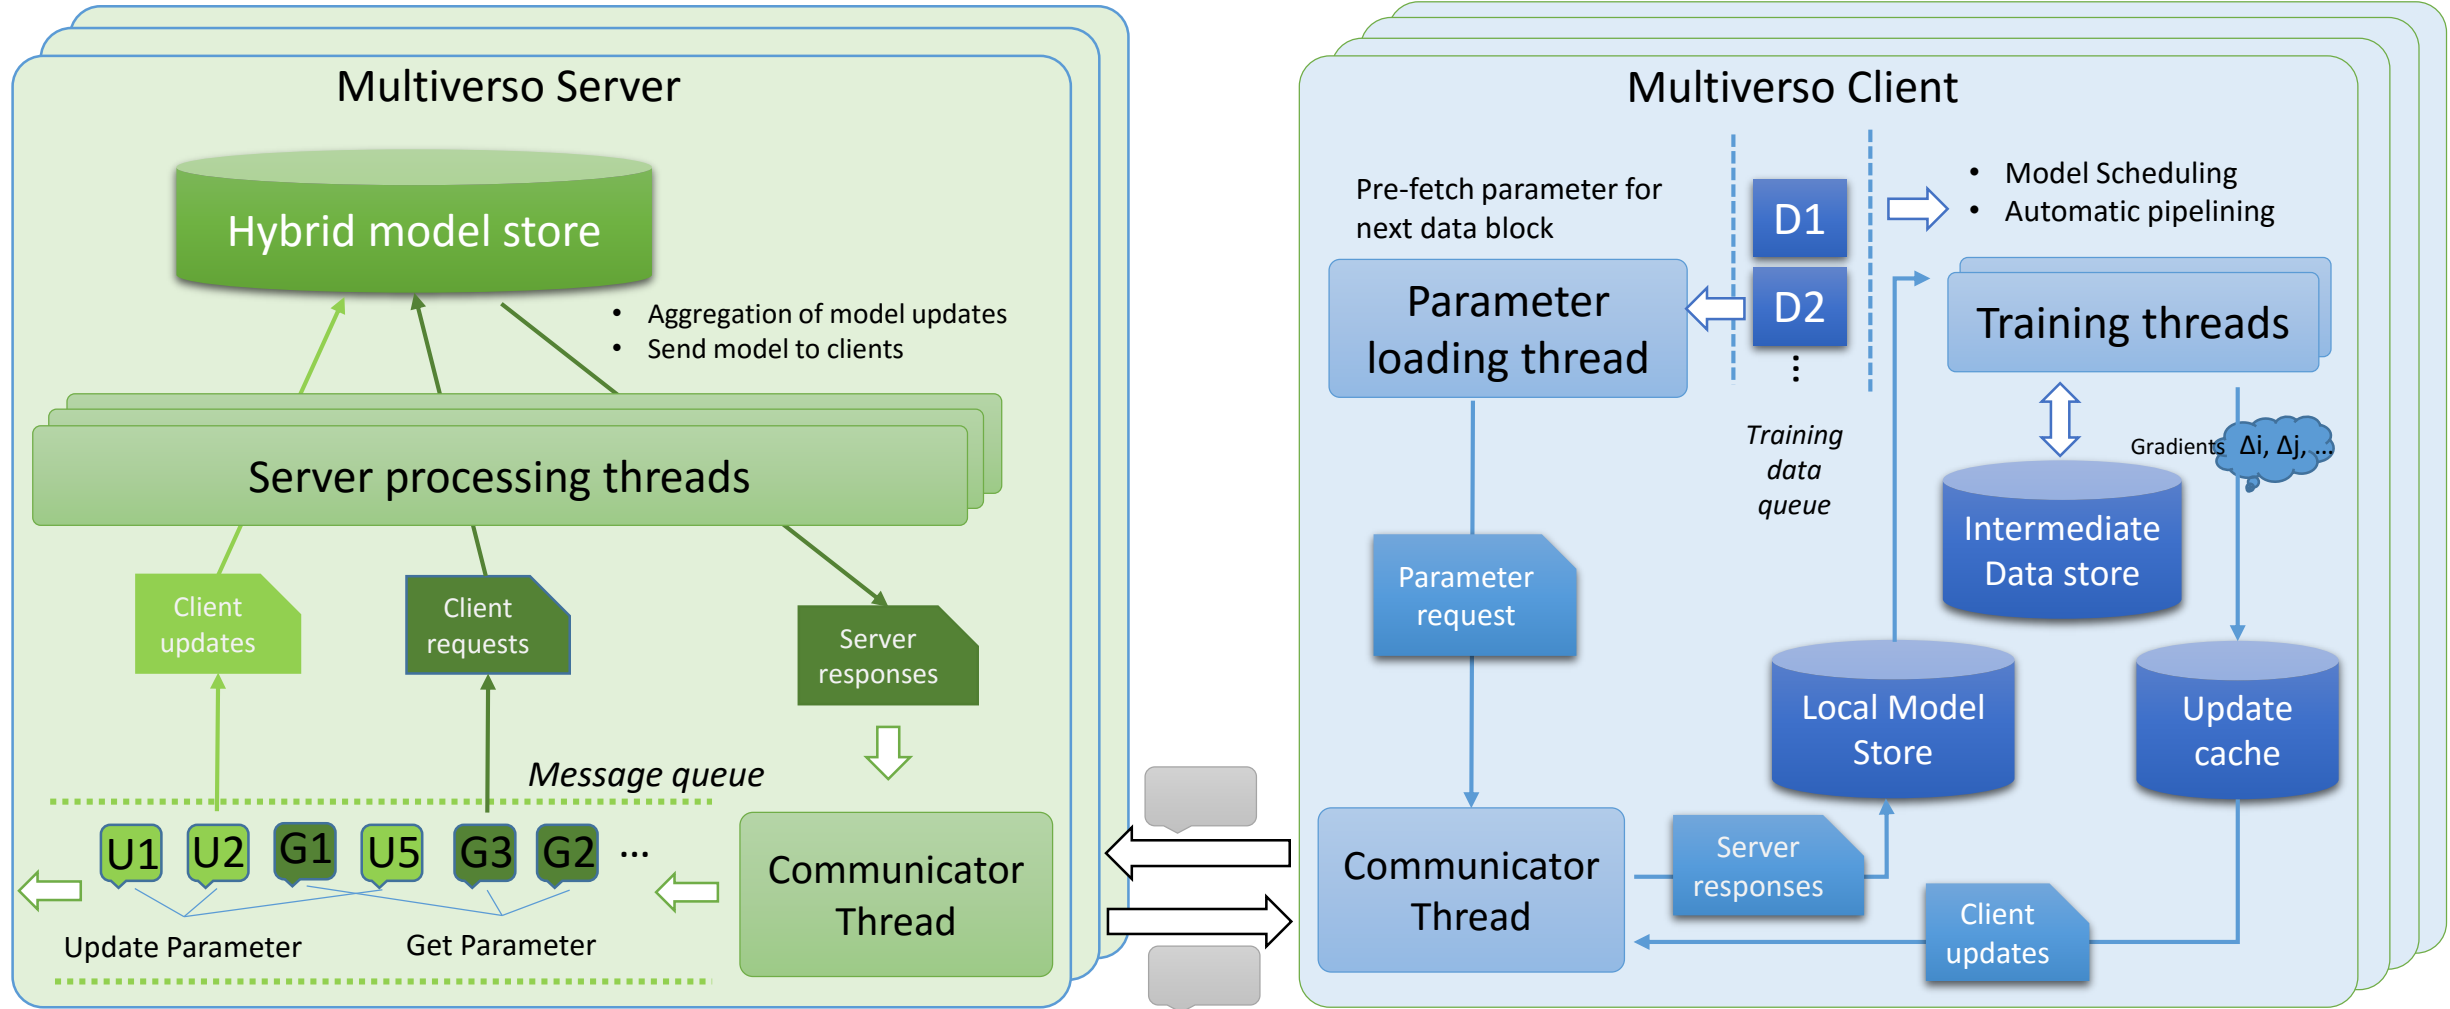


[1] Jinhui Yuan, Fei Gao, Qirong Ho, Wei Dai, Jinliang Wei, Xun Zheng, Eric Xing, Tie-Yan Liu, and Wei-Ying Ma, LightLDA: Topic Models on Modest Computer Cluster, WWW 2015.

[2] Tomas Mikolov, Kai Chen, Greg Corrado, and Jeffrey Dean. Efficient Estimation of Word Representations in Vector Space. ICLR 2013

[3] Fei Tian, Hanjun Dai, Jiang Bian, Bin Gao, Rui Zhang, Enhong Chen, and Tie-Yan Liu. A probabilistic model for learning multi-prototype word embeddings. COLING 2014.

Multiverso Framework



Software and Hardware Requirements

- Running platform - x64
 - Tested on Windows server 2012, and Ubuntu 12.04
- Developing environment and compilers
 - The codes were developed with Visual Studio 2013 on Windows; slight code conversion might be needed when using other versions of Visual Studio.
 - The codes were compiled with g++ 4.8 on Linux.
- Welcome to help us migrate DMTK to other environments

Using Binary Tools

Runing LightLDA and Distributed Word Embedding(DWE/DMWE)

Prerequisites

- Install MPI/ZMQ
 - <https://github.com/Microsoft/multiverso/tree/master/windows>
- Firewall exception
 - If your cluster has firewall enabled, please make sure you have firewall exception for the following network applications
 - All binaries for MPI, i.e., executables in the “bin” folder under the MPI path
 - The commands you are going to run on MPI/ZMQ, e.g., LightLDA.exe
 - Server applications (when using ZMQ as inter-process communication library)
- Make third-party dlls ready for applications
 - Mpi.dll, ZMQ-related dlls.
 - Make sure they are in the system path or in the application path

Running Multiverso

- Multiverso supports two inter-process communication modes
 - MPI:
 - The parameter server routines are bound to the worker processes running as server threads. MPI will help start all the applications.
 - Command line to start a MPI job: *mpiexec.exe -machinefile machineFile app.exe -yourAppArgs ...*
 - ZMQ:
 - The parameter server is a standalone process. You need to start servers and your client applications separately. For client applications, a server endpoint file is needed, which contains the *ip:port* information of each server line by line.
 - Go to every server and start the following commands:
 - To start server: *multiverso server.exe server id worker number*
 - To start a particular client task: *app.exe -server endpoint file server file -yourAppArgs ...*

More about MPI Jobs

- Follow the [Prerequisites page](#) to install MPI, and add firewall exceptions
- Before running the MPI job, please make sure you have started “SMPD.exe -p port_id -d” on each machine. The port_id is the port used to communicate among MPI nodes.
- Compose the machine.txt file to specify the machine names (or IPs) and node (or process) number on each machine.
- Run the following command to kickoff the distributed training on any machine within the cluster.

```
mpiexec -p port_id -machinefile machine.txt path\App.exe -arguments
```

MACHINE1 2
MACHINE2 4

- Remarks
 - Please make sure every machine has the working directory on their local disk.
 - Please make sure port_id is consistent with the value you start smpd.

In this example, machine1 and machine2 are used: 2 nodes (processes) on machine1, and 4 nodes (processes) on machine2.

Running LightLDA

- LightLDA only accepts binary input data; we have provided a tool to convert libsvm format text file to the desirable binary file.
- For more details, please refer to
 - <https://github.com/Microsoft/lightlda/tree/master/example>

Running DWE/DMWE

- DWE/DMWE only accepts raw text file as input. In addition, **Vocabulary file, stop words file** are also required. We have provided a tool to generate vocabulary file based on the raw text.
- For more details, please refer to
 - https://github.com/Microsoft/distributed_word_embedding/examples

Developing Your Own Algorithm

Build more new distributed machine learning algorithms based on Multiverso

Setup Development Environment

- Download

git clone https://github.com/Microsoft/multiverso

- Build

- On Linux

- Run `./third_party/install.sh`
 - Run `make all -j4`

- On Windows

- Install third-party libraries
<https://github.com/Microsoft/multiverso/tree/master/windows>
 - Open `windows/multiverso.sln`, change configuration and platform as Release and x64, then build the solution.

Key Steps

- Define the main training routine
 - 1) Start the Multiverso environment
 - 2) Configure the Multiverso parameter server
 - 3) Define the overall training process based on user defined data block
- Implement you algorithm logic
 - 1) Define your data block
 - 2) Define parameter loader
 - 3) Define training logic of each data block

Start the Multiverso Environment

- Initialize the Multiverso environment.
 - pass the configuration to the environment by setting a **Config** object.
 - call the **Init** function to start Multiverso.
 - Parameters of **Init** function include ***trainer*** and ***parameter loader***, which define the detailed training algorithm.
 - call **Close** to close Multiverso.
- **Remark**
 - More details could be found on [our website](#)
 - We also have comprehensive comments in the [source codes](#) to aid your programming

```
int main(int argc, char* argv[])
{
    // Step 1: Init Multiverso Environment
    // Set the configuration
    Config config;
    config.num_trainers = 1; config.num_servers = 1; config.max_delay = 0;

    // Trainer and Loader will be defined later
    std::vector<TrainerBase*> trainers(config.num_trainers, new Trainer());
    ParameterLoaderBase* loader = new ParameterLoader();

    Multiverso::Init(trainers, loader, config, &argc, &argv);

    // Step 2: Config the table in Parameter Server
    Multiverso::BeginConfig();

    const int kTableId = 0, kNumRows = 1, kNumCols = 10;
    // Create Table in Parameter Server
    Multiverso::AddTable(kTableId, kNumRows, kNumCols, Type::Int, Format::Dense);
    for (int k = 0; k < kNumCols; ++k)
        // Init the value in Parameter Server
        Multiverso::AddToServer<int>(kTableId, kNumRows, kNumCols, 1);

    // Finish Configuration
    Multiverso::EndConfig();

    // Step 3: Train
    const int kNumIteration = 100;
    Multiverso::BeginTrain();

    for (int i = 0; i < kNumIteration; ++i)
    {
        Multiverso::BeginClock();

        // DataBlock will be defined later
        DataBlock* data = GetDataBlock();
        Multiverso::PushDataBlock(data);

        Multiverso::EndClock();
    }
}
```

Configure Multiverso Parameter Server

- Multiverso stores parameters as **Table**. Before training, you need to configure and initialize the server tables by adding your logic between these two methods: **BeginConfig**, **EndConfig**.
- *Configure tables*
 - Call **AddTable** to create tables in both local cache and parameter server
 - Call **SetRow** method to set row property, like size, format, type, etc.
- *Initialize tables*
 - Call **AddToServer**, if you want to initialize the server tables with non-trivial values.

```
int main(int argc, char* argv[])
{
    // Step 1: Init Multiverso Environment
    // Set the configuration
    Config config;
    config.num_trainers = 1; config.num_servers = 1; config.max_delay = 0;

    // Trainer and Loader will be defined later
    std::vector<TrainerBase*> trainers(config.num_trainers, new Trainer());
    ParameterLoaderBase* loader = new ParameterLoader();

    Multiverso::Init(trainers, loader, config, &argc, &argv);

    // Step 2: Config the table in Parameter Server
    Multiverso::BeginConfig();

    const int kTableId = 0, kNumRows = 1, kNumCols = 10;
    // Create Table in Parameter Server
    Multiverso::AddTable(kTableId, kNumRows, kNumCols, Type::Int, Format::Dense);
    for (int k = 0; k < kNumCols; ++k)
        // Init the value in Parameter Server
        Multiverso::AddToServer<int>(kTableId, kNumRows, kNumCols, 1);

    // Finish Configuration
    Multiverso::EndConfig();

    // Step 3: Train
    const int kNumIteration = 100;
    Multiverso::BeginTrain();

    for (int i = 0; i < kNumIteration; ++i)
    {
        Multiverso::BeginClock();
```


Define Training Process

- The training logic should be placed between these two methods: **BeginTrain** and **EndTrain**. Then Multiverso will schedule the training process based on data blocks.
- In distributed training, the concept of ***clock*** is needed, which refers to a local training period. For each clock we do a sync-up with parameter server. To define a clock, you need to call **BeginClock** and **EndClock** to tell Multiverso that the data fed are within the clock.
- In a clock period, you may feed one or more data blocks into Multiverso by calling **PushDataBlock**.

```

Multiverso::Init(trainers, loader, config, &argc, &argv);

// Step 2: Config the table in Parameter Server
Multiverso::BeginConfig();

const int kTableId = 0, kNumRows = 1, kNumCols = 10;
// Create Table in Parameter Server
Multiverso::AddTable(kTableId, kNumRows, kNumCols, Type::Int, Format::Dense);
for (int k = 0; k < kNumCols; ++k)
    // Init the value in Parameter Server
    Multiverso::AddToServer<int>(kTableId, kNumRows, kNumCols, 1);

// Finish Configuration
Multiverso::EndConfig();

// Step 3: Train
const int kNumIteration = 100;
Multiverso::BeginTrain();

for (int i = 0; i < kNumIteration; ++i)
{
    Multiverso::BeginClock();

    // DataBlock will be defined later
    DataBlock* data = GetDataBlock();
    Multiverso::PushDataBlock(data);

    Multiverso::EndClock();
}

Multiverso::EndTrain();
// End of Multiverso: Close the Multiverso Environment
Multiverso::Close();
return 0;
}

```

Define Your Data Block

- Multiverso schedules the training process based on data blocks. To define your data block, you just need to inherit the **DataBlockBase** class and implement your own data block.
- In **DataBlockBase**, you might need to define the functions for
 - composing data block from input data stream (in the form of file or memory buffer).
 - getting training examples from data block.

```
class DataBlock : public DataBlockBase
{
    // Defines what's your training data
};

class Trainer : public TrainerBase
{
public:
    // Defines your training logic for a data block 'data'
    // When calling this function, you can assume that all data
    // and parameter needed is local. Parameters have been prefetched
    // from parameter server by multiverso. Then you can write your
    // training logic as same as writing a single machine program.
    void TrainIteration(DataBlockBase* data) override
    {
        DataBlock* data_block = reinterpret_cast<DataBlock*>(data);

        // The API GetRow<T> and Add<T> are member function inherited
        // from the TrainerBase

        // Access model, it is local memory access, no network happened here
        Row<int>& row = GetRow<int>(kTableId, kRowId);

        std::vector<int> updates;
        // Your training logic to produce updates
        // TODO
        // ...
        Add<int>(kTableId, kRowId, updates);
    }
};

class ParameterLoader : public ParameterLoaderBase
{

```

Define Parameter Loader

- You need to implement parameter preparation and training for each data block separately, so that Multiverso can pipeline the process to enhance system throughput.
- ParameterLoader** is used to prepare parameters needed by training. To implement your loader, you need to inherit the **ParameterLoaderBase** class and override the **ParseAndRequest** method. In the method, you need to
 - Implement logics to parse the data block and identify a set of parameters needed during the upcoming training process.
 - Use **RequestTable**, **RequestRow**, or **RequestElement** to pull these parameters from parameter servers.

```
void TrainIteration(DataBlockBase* data) override
{
    DataBlock* data_block = reinterpret_cast<DataBlock*>(data);

    // The API GetRow<T> and Add<T> are member function inherited
    // from the TrainerBase

    // Access model, it is local memory access, no network happened here
    Row<int>& row = GetRow<int>(kTableId, kRowId);

    std::vector<int> updates;
    // Your training logic to produce updates
    // TODO
    // ...
    Add<int>(kTableId, kRowId, updates);
}
};
```

```
class ParameterLoader : public ParameterLoaderBase
{
public:
    // Defines which parameter you need for training this data block
    void ParseAndRequest(DataBlockBase* data) override
    {
        RequestTable(kTableId);
        RequestRow(kTableId, kRowId);
    }
};
```

Define the Training Logic

- Multiverso will create training threads based on user-defined data blocks.
- These training threads will call **Trainer** to perform the training. To implement your trainer, please inherit the **TrainerBase** class and override the **TrainIteration** method.
- In the method, you can
 - implement your training logic to process the training samples in the data block
 - Use the **GetRow** method to access parameters in local cache
 - Use the **Add** method to update the parameters.

```
class DataBlock : public DataBlockBase
{
    // Defines what's your training data
};

class Trainer : public TrainerBase
{
public:
    // Defines your training logic for a data block 'data'
    // When calling this function, you can assume that all data
    // and parameter needed is local. Parameters have been prefetched
    // from parameter server by multiverso. Then you can write your
    // training logic as same as writing a single machine program.
    void TrainIteration(DataBlockBase* data) override
    {
        DataBlock* data_block = reinterpret_cast<DataBlock*>(data);

        // The API GetRow<T> and Add<T> are member function inherited
        // from the TrainerBase

        // Access model, it is local memory access, no network happened here
        Row<int>& row = GetRow<int>(kTableId, kRowId);

        std::vector<int> updates;
        // Your training logic to produce updates
        // TODO
        // ...
        Add<int>(kTableId, kRowId, updates);
    }
};
```

```
class ParameterLoader : public ParameterLoaderBase
```

Resources

- DMTK Website
<https://www.dmtk.io>
- DMTK Source Codes
<https://github.com/Microsoft/DMTK>
- DMTK Documents
<http://www.dmtk.io/document.html>
- Multiverso API Documents
<http://www.dmtk.io/multiverso/annotated.html>

Support

- Please send email to dmtk@microsoft.com for technical support or bug reporting.
- We will continue to enrich the tutorial to aid your development on top of DMTK; please check our website and GitHub site in a regular basis for new information.